

1. Datasets

The experiments use three synthetic collections of labelled images. The table below lists the number of classes and the size of the training and test split of each collection.

Table 1. Statistics of the three synthetic datasets.

Dataset	Classes	Train	Test
Alpha	10	50,000	10,000
Beta	100	73,257	26,032
Gamma	200	100,000	5,000

The text continues below the table with an ordinary paragraph, so that the table is surrounded by body text on both sides as it would be in a paper.

2. Results

Every method was trained three times. The table reports the mean accuracy, the mean F1 score and the training time in minutes, without any ruling lines between the cells.

Method	Accuracy	F1	Time
Linear	71.25	70.90	3.5
MLP	80.00	79.42	6.0
ConvNet	88.13	87.60	14.2
ResNet	92.10	91.40	31.0
ViT	94.70	94.10	58.4
Ours	95.32	95.01	40.7

Table 2. Comparison with baselines on the Alpha dataset.

The text continues below the table with an ordinary paragraph, so that the table is surrounded by body text on both sides as it would be in a paper.

3. Ablation

The last table groups two accuracy columns under one merged header cell. The size of every model is given in millions of parameters in the final column.

Table 3. Accuracy and size of the ablated models.

Model	Accuracy		Params
	Top-1	Top-5	
Small	81.20	95.10	5.3
Base	86.50	97.42	21.8
Large	88.10	98.00	86.0

The text continues below the table with an ordinary paragraph, so that the table is surrounded by body text on both sides as it would be in a paper.